

NOTE: All the Map reduce programs related to TASKS 1,2,3 are written and exported the programs as respective jar files. Programs are attached as Zipped files in Github.

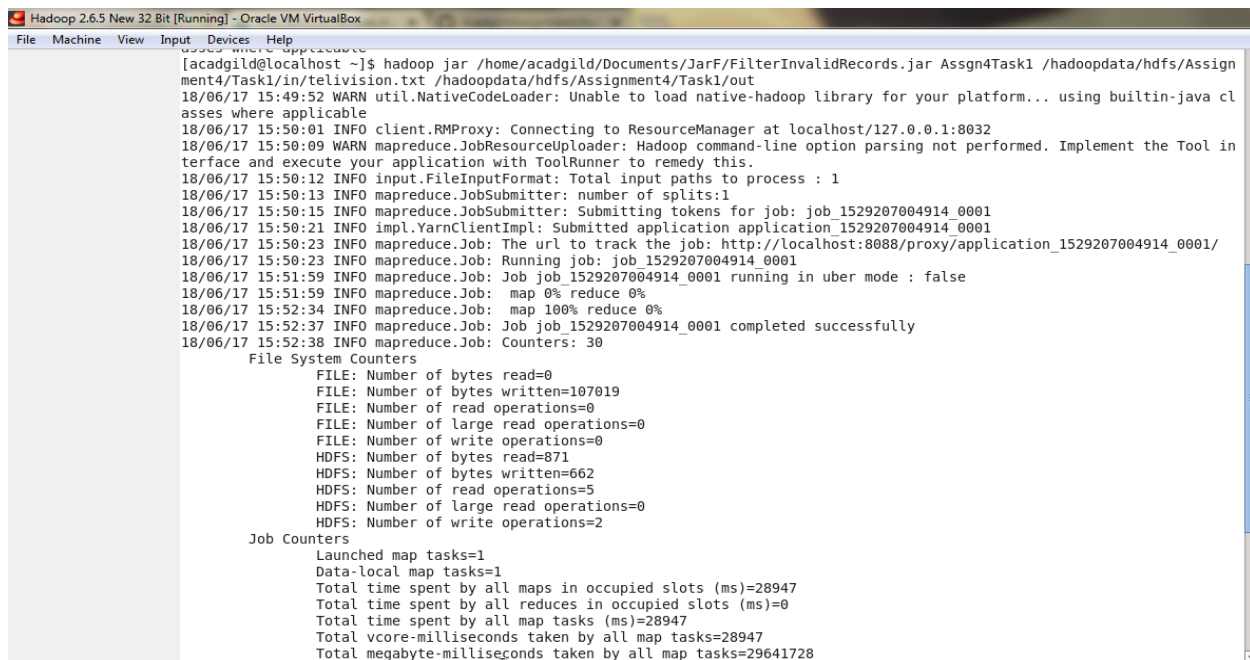
Task 1:

Write a Map Reduce program to filter out the invalid records. Map only job will fit for this context.

```
[acadgild@localhost ~]$ hadoop fs -mkdir /hadoopdata/hdfs/Assignment4/Task1/in
18/06/17 15:44:26 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
[acadgild@localhost ~]$ hadoop fs -put /home/acadgild/Downloads/television.txt /hadoopdata/hdfs/Assignment4/Task1/in
18/06/17 15:46:21 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
[acadgild@localhost ~]$ hadoop fs -mkdir /hadoopdata/hdfs/Assignment4/Task1/out
18/06/17 15:46:59 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

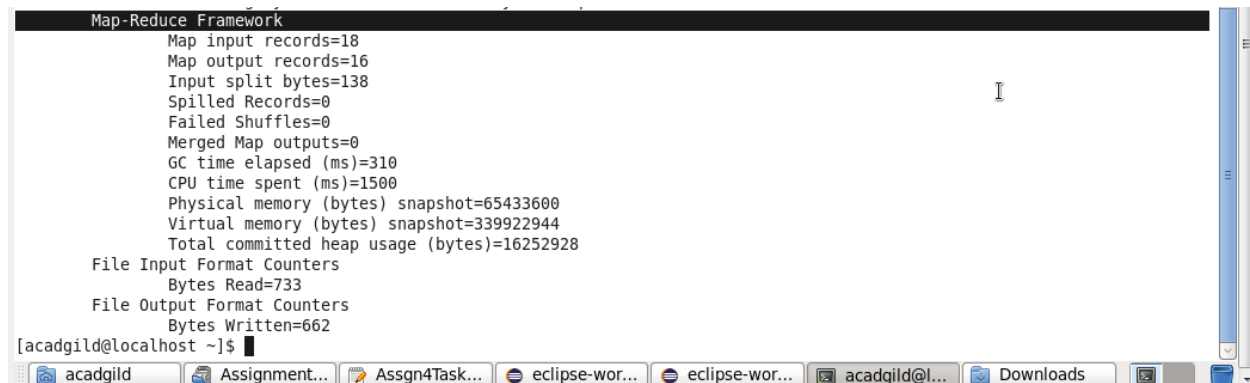
-----JAR FILE-----

```
[acadgild@localhost ~]$ hadoop jar /home/acadgild/Documents/JarF/FilterInvalidRecords.jar Assgn4Task1 /hadoopdata/hdfs/Assignment4/Task1/in/television.txt /hadoopdata/hdfs/Assignment4/Task1/out/
```



```
Hadoop 2.6.5 New 32 Bit [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help

[acadgild@localhost ~]$ hadoop jar /home/acadgild/Documents/JarF/FilterInvalidRecords.jar Assgn4Task1 /hadoopdata/hdfs/Assignment4/Task1/in/television.txt /hadoopdata/hdfs/Assignment4/Task1/out
18/06/17 15:49:52 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/06/17 15:50:01 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/06/17 15:50:09 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
18/06/17 15:50:12 INFO input.FileInputFormat: Total input paths to process : 1
18/06/17 15:50:13 INFO mapreduce.JobSubmitter: number of splits:1
18/06/17 15:50:15 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1529207004914_0001
18/06/17 15:50:21 INFO impl.YarnClientImpl: Submitted application application_1529207004914_0001
18/06/17 15:50:23 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1529207004914_0001/
18/06/17 15:50:23 INFO mapreduce.Job: Running job: job_1529207004914_0001
18/06/17 15:51:59 INFO mapreduce.Job: Job job_1529207004914_0001 running in uber mode : false
18/06/17 15:51:59 INFO mapreduce.Job: map 0% reduce 0%
18/06/17 15:52:34 INFO mapreduce.Job: map 100% reduce 0%
18/06/17 15:52:37 INFO mapreduce.Job: Job job_1529207004914_0001 completed successfully
18/06/17 15:52:38 INFO mapreduce.Job: Counters: 30
File System Counters
  FILE: Number of bytes read=0
  FILE: Number of bytes written=107019
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=871
  HDFS: Number of bytes written=662
  HDFS: Number of read operations=5
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
Job Counters
  Launched map tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=28947
  Total time spent by all reduces in occupied slots (ms)=0
  Total time spent by all map tasks (ms)=28947
  Total vcore-milliseconds taken by all map tasks=28947
  Total megabyte-milliseconds taken by all map tasks=29641728
```



```
Map-Reduce Framework
  Map input records=18
  Map output records=16
  Input split bytes=138
  Spilled Records=0
  Failed Shuffles=0
  Merged Map outputs=0
  GC time elapsed (ms)=310
  CPU time spent (ms)=1500
  Physical memory (bytes) snapshot=65433600
  Virtual memory (bytes) snapshot=339922944
  Total committed heap usage (bytes)=16252928
File Input Format Counters
  Bytes Read=733
File Output Format Counters
  Bytes Written=662
[acadgild@localhost ~]$
```

OUTPUT: output consists of records which do not have NA.

```
Bytes Written=662
[acadgild@localhost ~]$ hadoop fs -ls /hadoopdata/hdfs/Assignment4/Task1/out
18/06/17 16:09:48 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java cl
asses where applicable
Found 2 items
-rw-r--r-- 1 acadgild supergroup 0 2018-06-17 15:52 /hadoopdata/hdfs/Assignment4/Task1/out/_SUCCESS
-rw-r--r-- 1 acadgild supergroup 662 2018-06-17 15:52 /hadoopdata/hdfs/Assignment4/Task1/out/part-m-000000
[acadgild@localhost ~]$ hadoop fs -cat /hadoopdata/hdfs/Assignment4/Task1/out/part-m-000000
18/06/17 16:10:18 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java cl
asses where applicable
Samsung|Optima|14|Madhya Pradesh|132401|14200
Onida|Lucid|18|Uttar Pradesh|232401|16200
Akai|Decent|16|Kerala|922401|12200
Lava|Attention|20|Assam|454601|24200
Zen|Super|14|Maharashtra|619082|9200
Samsung|Optima|14|Madhya Pradesh|132401|14200
Onida|Lucid|18|Uttar Pradesh|232401|16200
Onida|Decent|14|Uttar Pradesh|232401|16200
Lava|Attention|20|Assam|454601|24200
Zen|Super|14|Maharashtra|619082|9200
Samsung|Optima|14|Madhya Pradesh|132401|14200
Samsung|Decent|16|Kerala|922401|12200
Lava|Attention|20|Assam|454601|24200
Samsung|Super|14|Maharashtra|619082|9200
Samsung|Super|14|Maharashtra|619082|9200
Samsung|Super|14|Maharashtra|619082|9200
```

TASK 2 : Write a Map Reduce program to calculate the total units sold for each Company.

-----JAR FILE-----

```
[acadgild@localhost ~]$ hadoop jar /home/acadgild/Documents/JarF/Task2.jar com.acadgild.Assgn4Task2.Assgn4Task2 /hadoopdata/
hdfs/Assignment4/Task1/in/television.txt /hadoopdata/hdfs/Assignment4/Task1/out
```

```
[Running] - Oracle VM VirtualBox
Input Devices Help
Samsung|Super|14|Maharashtra|619082|9200
[acadgild@localhost ~]$ hadoop jar /home/acadgild/Documents/JarF/FilterInvalidRecords.jar Task1Mapper /hadoopdata/hdfs/Assign
ment4/Task1/in/television.txt /hadoopdata/hdfs/Assignment4/Task1/out
Exception in thread "main" java.lang.NoSuchMethodException: Task1Mapper.main([Ljava.lang.String;)
    at java.lang.Class.getMethod(Class.java:1786)
    at org.apache.hadoop.util.RunJar.run(RunJar.java:215)
    at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
[acadgild@localhost ~]$ hadoop jar /home/acadgild/Documents/JarF/Task2.jar com.acadgild.Assgn4Task2.Assgn4Task2 /hadoopdata/
hdfs/Assignment4/Task1/in/television.txt /hadoopdata/hdfs/Assignment4/Task1/out
18/06/17 16:53:49 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java cl
asses where applicable
18/06/17 16:53:58 INFO client.RMPProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/06/17 16:54:03 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool in
terface and execute your application with ToolRunner to remedy this.
18/06/17 16:54:05 INFO input.FileInputFormat: Total input paths to process : 1
18/06/17 16:54:06 INFO mapreduce.JobSubmitter: number of splits:1
18/06/17 16:54:07 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1529207004914_0002
18/06/17 16:54:09 INFO impl.YarnClientImpl: Submitted application application_1529207004914_0002
18/06/17 16:54:10 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1529207004914_0002/
18/06/17 16:54:10 INFO mapreduce.Job: Running job: job_1529207004914_0002
18/06/17 16:54:52 INFO mapreduce.Job: Job job_1529207004914_0002 running in uber mode : false
18/06/17 16:54:52 INFO mapreduce.Job: map 0% reduce 0%
18/06/17 16:55:25 INFO mapreduce.Job: map 100% reduce 0%
18/06/17 16:56:02 INFO mapreduce.Job: map 100% reduce 100%
18/06/17 16:56:07 INFO mapreduce.Job: Job job_1529207004914_0002 completed successfully
18/06/17 16:56:08 INFO mapreduce.Job: Counters: 49
File System Counters
  FILE: Number of bytes read=225
  FILE: Number of bytes written=215831
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=871
  HDFS: Number of bytes written=43
  HDFS: Number of read operations=6
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
```

acadgild Assignment4.zip Task2Reducer.jar eclipse-worksp... eclipse-worksp... acadgild@local...

```
[Running] - Oracle VM VirtualBox
put Devices Help
Total megabyte-milliseconds taken by all reduce tasks=37696512
Map-Reduce Framework
  Map input records=18
  Map output records=18
  Map output bytes=183
  Map output materialized bytes=225
  Input split bytes=138
  Combine input records=0
  Combine output records=0
  Reduce input groups=6
  Reduce shuffle bytes=225
  Reduce input records=18
  Reduce output records=6
  Spilled Records=36
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=983
  CPU time spent (ms)=4000
  Physical memory (bytes) snapshot=236191744
  Virtual memory (bytes) snapshot=680849408
  Total committed heap usage (bytes)=139776000
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=733
File Output Format Counters
  Bytes Written=43
```

OUTPUT : output consists of total units sold for each Company

```
bytes written=733
[acadgild@localhost ~]$ hadoop fs -ls /hadoopdata/hdfs/Assignment4/Task1/out
18/06/17 16:57:50 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java cl
asses where applicable
Found 2 items
-rw-r--r-- 1 acadgild supergroup 0 2018-06-17 16:56 /hadoopdata/hdfs/Assignment4/Task1/out/_SUCCESS
-rw-r--r-- 1 acadgild supergroup 43 2018-06-17 16:56 /hadoopdata/hdfs/Assignment4/Task1/out/part-r-000000
[acadgild@localhost ~]$ hadoop fs -cat /hadoopdata/hdfs/Assignment4/Task1/out/part-r-000000
18/06/17 16:58:23 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java cl
asses where applicable
Akai 1
Lava 3
NA 1
Onida 4
Samsung 7
Zen 2
[acadgild@localhost ~]$
```

TASK 3: Write a Map Reduce program to calculate the total units sold in each state for Onida company.

---JAR FILE---

```
[acadgild@localhost ~]$ hadoop jar /home/acadgild/Documents/JarF/Task3.jar com.acadgild.Assgn4Task3.Assgn4Task3 /hadoopdata/hdfs/Assignment4/Task1/in/television.txt /hadoopdata/hdfs/Assignment4/Task1/out
```

```
[acadgild@localhost ~]$ hadoop jar /home/acadgild/Documents/JarF/Task3.jar com.acadgild.Assgn4Task3.Assgn4Task3 /hadoopdata/hdfs/Assignment4/Task1/in/television.txt /hadoopdata/hdfs/Assignment4/Task1/out
18/06/17 17:10:05 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/06/17 17:10:12 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/06/17 17:10:17 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
18/06/17 17:10:19 INFO input.FileInputFormat: Total input paths to process : 1
18/06/17 17:10:19 INFO mapreduce.JobSubmitter: number of splits:1
18/06/17 17:10:21 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1529207004914_0003
18/06/17 17:10:22 INFO impl.YarnClientImpl: Submitted application application_1529207004914_0003
18/06/17 17:10:23 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1529207004914_0003/
18/06/17 17:10:23 INFO mapreduce.Job: Running job: job_1529207004914_0003
18/06/17 17:11:07 INFO mapreduce.Job: Job job_1529207004914_0003 running in uber mode : false
18/06/17 17:11:07 INFO mapreduce.Job: map 0% reduce 0%
18/06/17 17:11:38 INFO mapreduce.Job: map 100% reduce 0%
18/06/17 17:12:10 INFO mapreduce.Job: map 100% reduce 100%
18/06/17 17:12:12 INFO mapreduce.Job: Job job_1529207004914_0003 completed successfully
18/06/17 17:12:13 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=79
    FILE: Number of bytes written=215539
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=871
    HDFS: Number of bytes written=41
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=27215
    Total time spent by all reduces in occupied slots (ms)=28067
```

it [Running] - Oracle VM VirtualBox

```
Input  Devices  Help
Total vcore-milliseconds taken by all reduce tasks=28067
Total megabyte-milliseconds taken by all map tasks=27868160
Total megabyte-milliseconds taken by all reduce tasks=28740608
Map-Reduce Framework
  Map input records=18
  Map output records=4
  Map output bytes=65
  Map output materialized bytes=79
  Input split bytes=138
  Combine input records=0
  Combine output records=0
  Reduce input groups=2
  Reduce shuffle bytes=79
  Reduce input records=4
  Reduce output records=2
  Spilled Records=8
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=1016
  CPU time spent (ms)=3670
  Physical memory (bytes) snapshot=236302336
  Virtual memory (bytes) snapshot=680697856
  Total committed heap usage (bytes)=139730944
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=733
File Output Format Counters
  Bytes Written=41
```

OUTPUT: Output consists of total units sold in each state for Onida company.

```
bytes written=41
[acadgild@localhost ~]$ hadoop fs -ls /hadoopdata/hdfs/Assignment4/Task1/out
18/06/17 17:14:50 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java cl
asses where applicable
Found 2 items
-rw-r--r--  1 acadgild supergroup          0 2018-06-17 17:12 /hadoopdata/hdfs/Assignment4/Task1/out/ SUCCESS
-rw-r--r--  1 acadgild supergroup        41 2018-06-17 17:12 /hadoopdata/hdfs/Assignment4/Task1/out/part-r-00000
[acadgild@localhost ~]$ hadoop fs -cat /hadoopdata/hdfs/Assignment4/Task1/out/part-r-00000
18/06/17 17:18:02 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java cl
asses where applicable
Onida : Kerala 1
Onida : Uttar Pradesh 3
[acadgild@localhost ~]$
```